

Performance Optimization and Profiling of MLPCpp in SU2

GSoC 2025 Final Report

Divyaprakash

August 31, 2025

1 Introduction

SU2 [1] is an open-source suite for computational fluid dynamics (CFD) widely used in both research and industrial design. Within SU2, the **MLPCpp** module provides a compact C++ library for enabling multilayer perceptron (MLP) inference as part of solver workflows. This capability supports a range of data-driven tasks, including the modeling of fluid properties.

The original implementation of **MLPCpp** contained some performance bottlenecks and limitations in usability for advanced data-driven workflows. As part of this Google Summer of Code (GSoC) project, a series of profiling, optimization, and refactoring efforts were undertaken to improve runtime efficiency and extend the functionality of the library. Key directions included the introduction of more efficient data structures, the refactoring of computational kernels for improved memory locality and numerical precision, and enhancements to variable mapping mechanisms to support more complex workflows.

Profiling was initially conducted with **gprof**, followed by more detailed analysis with **valgrind** and **kcachegrind**, and later **Tracy**, to identify bottlenecks in the neural network inference pipeline. Guided by these insights, multiple optimizations were implemented, ranging from memory layout improvements in the weight matrices to function-level optimizations using fused multiply-add instructions. In addition, a low-overhead profiling capability was integrated into the SU2 build system along with documentation to enable future contributors to easily instrument and profile the code.

2 Code Profiling and Identifying Bottlenecks

Code profiling was initially carried out using the **gprof** profiler that is included with the SU2 installation, in order to identify the functions contributing most to the overall runtime. However, to obtain a more complete picture of the child processes and to determine which operations were consuming time within those functions, **valgrind** and **kcachegrind** were subsequently employed. The latter provided an interactive visualization of the profiling data, including call graphs and annotated source code views, as shown in Figure 1. A detailed description of this process was also documented in a blog post.¹

From the profiling results, it was observed that the function **Predict** itself did not consume a significant amount of self time. Instead, the majority of the runtime was spent in the functions **ComputePsi**, **ComputeChi**, and **ComputeX**. These functions were therefore identified as the main targets for optimization.

2.1 Suggested Improvements

The performance analysis indicated that the loop within **ComputeX** could be made more efficient. A representative code snippet is shown below:

```
1 mlpdouble ComputeX(std::size_t iLayer, std::size_t iNeuron) const {
2     mlpdouble x;
3     x = total_layers[iLayer]->GetBias(iNeuron);
```

¹<https://dpcfd.com/posts/2025/06/gsoc-week-2/>




Figure 1: Call graph generated using valgrind and kcache-grind. The functions ComputePsi, ComputeChi, and ComputeX were identified as major contributors to the runtime.

```

4  std::size_t nNeurons_previous = total_layers[iLayer - 1]->GetNNeurons();
5  for (std::size_t jNeuron = 0; jNeuron < nNeurons_previous; jNeuron++) {
6      x += weights_mat[iLayer - 1][iNeuron][jNeuron] *
7          total_layers[iLayer - 1]->GetOutput(jNeuron);
8  }
9  return x;
10 }
```

The nested access pattern, arising from the use of a *vector of vectors of vectors*, was found to result in fragmented memory allocation:

```

1  std::vector<std::vector<std::vector<mlpdouble>>> weights_mat;
```

To improve performance, this structure was modified, with the changes and their effects described in the following sections.

2.2 Advanced Profiling with Tracy

In addition to valgrind, another profiling option was later explored, namely Tracy, which offers very low overhead and provides fine-grained insights into performance. This profiler was integrated with the meson build system of SU2, and changes were made to the SU2 codebase to facilitate easy instrumentation. Both the integration and the accompanying documentation were contributed as two separate Pull Requests as part of this GSoC project.²³

3 Optimizing Weights Matrix Storage

From the above discussions, it has been well established that one way to optimize the code is by improving how the weight matrices are stored. Currently, they are implemented as 3D nested vectors. To enhance performance, we consider two alternative approaches: replacing the 3D nested vectors with raw 3D pointers,

²<https://github.com/su2code/SU2/pull/2536>

³<https://github.com/su2code/su2code.github.io/pull/180>

or using flattened vectors. The implementation and results of both approaches are presented in the following subsections. Before that, however, we first review the current implementation of the weight matrix.

3.1 Weights Matrix

Before modifying the implementation of the weights matrix, it is important to understand the existing design. Neural networks stored in `.mlp` files follow the TensorFlow convention, where weights are arranged as *source* $\times$ *destination*. In contrast, traditional mathematical notation typically uses the *destination* $\times$ *source* order. This difference is illustrated in Figure 2.

The `MLP` C++ code addresses this discrepancy in a clean manner:

1. `CReadNeuralNetwork::weights_mat` serves as a temporary container inside the reader class. It parses the `.mlp` file and stores the weights exactly as they appear.
2. `CNeuralNetwork::weights_mat` belongs to the functional network class. Here, the matrix is transposed to match the conventional layout and is used in the `Predict` method for evaluation.

The key reason for this layout is computational efficiency. CPUs favor memory locality: they are significantly faster when reading data stored contiguously. Consider the dot product inside the `ComputeX` function:

```

1  // Core computation of the network
2  for (std::size_t jNeuron = 0; jNeuron < nNeurons_previous; jNeuron++) {
3      x += weights_mat[iLayer - 1][iNeuron][jNeuron] *
4          total_layers[iLayer - 1]->GetOutput(jNeuron);
5  }
```

Here, the input for a destination neuron (`iNeuron`) is computed by iterating through all source neurons (`jNeuron`). With the current `[layer][destination][source]` storage, all weights for a given destination neuron are contiguous in memory. During iteration, the CPU accesses them sequentially.

From the above discussions, it is concluded that any modification to the weight matrix storage for performance benefits must be applied to `CNeuralNetwork::weights_mat`.

3.2 3D Pointer

The neural network implementation previously employed a three-dimensional vector to represent the weight matrix.

```
std::vector<std::vector<std::vector<mlpdouble>>> weights_mat
```

While convenient, this approach introduced additional overhead and less predictable memory access patterns. To optimize performance and provide finer control over memory allocation, the structure has been refactored to use a triple pointer, declared as

```
mlpdouble*** weights_mat
```

3.2.1 Implementation

The refactored design maintains the logical indexing scheme

```
weights_mat[layer][destination][source]
```

while eliminating the overhead of nested `std::vector` containers. Allocation proceeds hierarchically:

1. An array of pointers for each layer.
2. For each layer, an array of neuron pointers.
3. For each neuron, an array of incoming weights.




Figure 2: Comparison of weight matrix arrangements in neural networks: conventional mathematical notation (a) versus TensorFlow implementation (b). The diagrams show connections between a 2-neuron source layer and a 3-neuron destination layer, with corresponding weight matrix dimensions and indexing conventions.

It is implemented as follows in the `SizeWeights()` Function.

```

1  weights_mat = new mlpdouble*[n_hidden_layers + 1];
2  weights_mat[0] = new mlpdouble*[hiddenLayers[0]->GetNNNeurons()];
3  for (auto iNeuron = 0u; iNeuron < hiddenLayers[0]->GetNNNeurons(); iNeuron++)
4      weights_mat[0][iNeuron] = new mlpdouble[inputLayer->GetNNNeurons()];
5
6  for (auto iLayer = 1u; iLayer < n_hidden_layers; iLayer++) {
7      weights_mat[iLayer] = new mlpdouble*[hiddenLayers[iLayer]->GetNNNeurons()];
8      for (auto iNeuron = 0u; iNeuron < hiddenLayers[iLayer]->GetNNNeurons();
9          iNeuron++) {
10         weights_mat[iLayer][iNeuron] = new mlpdouble[hiddenLayers[iLayer] -
11             ↪ 1]->GetNNNeurons()];
12     }
13 }
14 weights_mat[n_hidden_layers] = new mlpdouble*[outputLayer->GetNNNeurons()];
15 for (auto iNeuron = 0u; iNeuron < outputLayer->GetNNNeurons(); iNeuron++) {
16     weights_mat[n_hidden_layers][iNeuron] = new mlpdouble[hiddenLayers[n_hidden_layers
17         ↪ - 1]->GetNNNeurons()];
18 }
19
20 ANN_outputs = new mlpdouble[outputLayer->GetNNNeurons()];

```

Because memory is allocated on the heap, manual deallocation must follow the following order:

1. Deallocate each neuron's weight arrays (which store the connection weights from the previous layer)
2. Deallocate each layer's array of pointers to neurons.
3. Finally, deallocate the outer layer array.

This approach ensures memory safety and prevents leaks. The updated destructor is given below:

```

1  ~CNeuralNetwork() {
2      delete[] ANN_outputs;
3
4      if (weights_mat != nullptr) {
5          // Clean up first layer (input to first hidden)
6          if (n_hidden_layers > 0) {
7              for (auto iNeuron = 0u; iNeuron < hiddenLayers[0]->GetNNeurons(); iNeuron++)
8                  ↪ {
9                  delete[] weights_mat[0][iNeuron];
10             }
11             delete[] weights_mat[0];
12         }
13
14         // Clean up hidden layers
15         for (auto iLayer = 1u; iLayer < n_hidden_layers; iLayer++) {
16             for (auto iNeuron = 0u; iNeuron < hiddenLayers[iLayer]->GetNNeurons();
17                 ↪ iNeuron++) {
18                 delete[] weights_mat[iLayer][iNeuron];
19             }
20             delete[] weights_mat[iLayer];
21         }
22
23         // Clean up last layer (last hidden to output)
24         if (outputLayer != nullptr) {
25             for (auto iNeuron = 0u; iNeuron < outputLayer->GetNNeurons(); iNeuron++) {
26                 delete[] weights_mat[n_hidden_layers][iNeuron];
27             }
28             delete[] weights_mat[n_hidden_layers];
29         }
30
31         delete[] weights_mat;
32     }
33
34     delete inputLayer;
35     delete outputLayer;
36     for (std::size_t i = 1; i < total_layers.size() - 1; i++) {
37         delete total_layers[i];
38     }
39 };

```

The conversion from `std::vector` to a triple-pointer preserves the computational logic while improving control over allocation. Public interfaces remain unchanged, so expressions such as `weights_mat[i][j][k]` continue to function as before. The main trade-off is between dynamic flexibility and simplicity: the vector-based implementation offered dynamic memory allocation capabilities (such as adding or removing elements at runtime), whereas the pointer-based design uses a simpler, more direct approach. Since the weights need to be sized only once and do not require runtime resizing, the dynamic allocation features that vectors provide are unnecessary for this application. Following the KISS (keep it simple, stupid) principle, the pointer-based approach selects the lightest, simplest tool that fulfills all requirements, potentially providing performance

benefits due to reduced overhead and improved cache locality.

Implementation Note: The complete source code for this refactoring approach is available in the `gsoc_dp` branch of the project repository⁴.

3.3 Flattened Vector

Disclaimer

Some of the changes described in this section are still under development and may not yet work as intended.

In this section, an alternative approach is proposed to refactor the core weight storage from a 3D `std::vector` of vectors to a single, contiguous, flattened 1D `std::vector` for improved performance. The refactoring is to be performed in three safe, verifiable phases:

1. **Parallel Implementation:** Introduce the new flattened vector and its accessor functions *alongside* the existing implementation.
2. **Verification:** Run both the original `Predict` function and a new `Predict_flat` function to check if the new implementation is mathematically identical.
3. **Cleanup:** Once verification is successful, the old 3D vector and its associated functions are removed.

3.3.1 Implementation

This phase involves adding new code without modifying the existing, validated logic.

Step 1: Modify `CNeuralNetwork.hpp` - Data Structures & Accessors: New data structures and accessor functions are added to the `CNeuralNetwork` class. In the `private` section of the `CNeuralNetwork` class, new variables are added:

```
1 // In the private section of CNeuralNetwork
2 // --- New flattened vector implementation ---
3 std::vector<mlpdouble> flat_weights;
4 std::vector<int> layer_sizes;
5 std::vector<int> layer_offsets;
```

New private and public helper functions are added below. These functions utilize the `inline` keyword to ensure compiler optimization eliminates function call overhead when optimizations are enabled (e.g., with the `-O2` flag).

```
1 // In the private section of CNeuralNetwork
2 /**
3  * @brief (NEW) Calculates the 1D index for the flattened weights vector.
4  */
5 inline int get_flat_index(int i_layer, int i_dest_neuron, int i_src_neuron) const {
6     return layer_offsets[i_layer] + i_dest_neuron * layer_sizes[i_layer] + i_src_neuron;
7 }
8
9 // In the public section of CNeuralNetwork
10 /**
```

⁴https://github.com/divyaprakash-iitd/MLPCpp/tree/gsoc_dp

```

11  * @brief (NEW) Sets a weight value in the flattened vector, handling the index swap.
12  */
13  inline void SetFlatWeight(unsigned long i_layer, unsigned long i_src_neuron,
14                          unsigned long i_dest_neuron, mlpdouble value) {
15      int index = get_flat_index(i_layer, i_dest_neuron, i_src_neuron);
16      flat_weights[index] = value;
17  }
18
19  /**
20   * @brief (NEW) Gets a weight value from the flattened vector using the in-memory
21   * ↪ layout.
22   */
23  inline mlpdouble GetFlatWeight(int i_layer, int i_dest_neuron, int i_src_neuron) const {
24      return flat_weights[get_flat_index(i_layer, i_dest_neuron, i_src_neuron)];
25  }

```

Step 2: Modify CNeuralNetwork.hpp - Sizing the New Structures: In the SizeWeights() function, logic is added to calculate the total number of weights and correctly size the new helper vectors and the flat_weights vector.

```

1  // File: CNeuralNetwork.hpp
2  void SizeWeights() {
3      // ... existing code to set up total_layers ...
4
5      // --- START NEW IMPLEMENTATION ---
6      layer_sizes.clear();
7      for(const auto& layer : total_layers) {
8          layer_sizes.push_back(layer->GetNNeurons());
9      }
10
11      layer_offsets.clear();
12      int total_weights = 0;
13      layer_offsets.push_back(total_weights);
14
15      for (size_t i = 0; i < GetNLayers() - 1; ++i) {
16          int weights_in_layer = layer_sizes[i+1] * layer_sizes[i];
17          total_weights += weights_in_layer;
18          layer_offsets.push_back(total_weights);
19      }
20      flat_weights.resize(total_weights);
21      // --- END NEW IMPLEMENTATION ---
22
23      // --- OLD IMPLEMENTATION (unchanged for now) ---
24      weights_mat.resize(n_hidden_layers + 1);
25      // ... (rest of the old implementation)
26  }

```

Step 3: Modify CLookUp_ANN.hpp - Populating Both Data Structures: This represents the only change needed for the data loading part of the refactoring. In the CLookUp_ANN class, within the nested loop that sets the weights, a call to the new SetFlatWeight function is added. The Reader object has already parsed the .ann file into its own memory. This loop transfers that data into the ANN object used for computations. A line is simply added to populate the new flat_weights vector simultaneously with the original weights_mat.

```

1 // File: CLookUp_ANN.hpp
2 for (auto i_layer = 0u; i_layer < ANN.GetNWeightLayers(); i_layer++) {
3     ANN.SetActivationFunction(i_layer, Reader.GetActivationFunction(i_layer));
4     for (auto i_neuron = 0u; i_neuron < ANN.GetNNeurons(i_layer); i_neuron++) {
5         for (auto j_neuron = 0u; j_neuron < ANN.GetNNeurons(i_layer + 1); j_neuron++) {
6
7             // Original call (unchanged)
8             ANN.SetWeight(i_layer, i_neuron, j_neuron,
9                 Reader.GetWeight(i_layer, i_neuron, j_neuron));
10
11             // New call to populate the parallel flat vector structure
12             ANN.SetFlatWeight(i_layer, i_neuron, j_neuron,
13                 Reader.GetWeight(i_layer, i_neuron, j_neuron));
14             // -----
15         }
16     }
17 }

```

Step 4: Modify CNeuralNetwork.hpp - Create Parallel Compute Functions: For every function that uses weights_mat, a duplicate version with a _flat suffix is created that uses the new GetFlatWeight accessor.

```

1 // File: CNeuralNetwork.hpp
2
3 /**
4  * @brief (NEW) Compute neuron activation function input using the flattened weights.
5  */
6 mlpdouble ComputeX_flat(std::size_t iLayer, std::size_t iNeuron) const {
7     mlpdouble x = total_layers[iLayer]->GetBias(iNeuron);
8     std::size_t nNeurons_previous = total_layers[iLayer - 1]->GetNNeurons();
9     for (std::size_t jNeuron = 0; jNeuron < nNeurons_previous; jNeuron++) {
10         // Access is [layer][destination][source]
11         x += GetFlatWeight(iLayer - 1, iNeuron, jNeuron) *
12             total_layers[iLayer - 1]->GetOutput(jNeuron);
13     }
14     return x;
15 }
16
17 // ... and so on for ComputePsi, ComputeChi, etc. ...

```

Once verification with the flattened implementations is successful, the old implementation can be safely removed. After these steps, the project is fully migrated to the new, more efficient flattened weight structure, providing improved memory locality and cache performance for neural network computations.

Implementation Note: The complete source code for this refactoring approach is available in the `flatten` branch of the project repository⁵.

4 Optimizing Variable Matching

Efficient mapping of input and output variables is essential for integrating neural networks within MLPCpp, particularly when dealing with multiple networks that have different input and output configurations. In the combustion tutorial case, for example, there are multiple networks with different outputs, requiring flexible variable mapping to correctly interface with the simulation code. A hash map-based method, along with a cleaner replacement for the `FindVariableIndices` function, was proposed to improve clarity and maintainability.

4.1 Hash Maps

The changes were made to the `FindVariableIndices` function. This function is responsible for mapping variable names to their corresponding indices within the Artificial Neural Network (ANN).

The original implementation used a nested loop structure to find matching variable names. The outer loop iterated through the variables of the ANN, and the inner loop iterated through the list of variable names to be found. This approach has a time complexity of $O(N \times M)$, where N is the number of variables in the ANN and M is the number of variables to be mapped.

4.1.1 Implementation

The new implementation in `CLookUp_ANN.hpp` replaces the nested loops with a more efficient approach utilizing a `std::unordered_map` (hash map). The `std::unordered_map` provides $O(1)$ average-case lookup time vs `std::map`'s $O(\log n)$, which is crucial when mapping variable names since we're doing many lookups and don't need the keys to be sorted. The implementation steps are as follows:

1. A hash map is created, mapping each variable name to its index in the input list.
2. The code then iterates a single time through the ANN's variables.
3. For each ANN variable, it performs a quick lookup in the hash map to see if there is a match.

This change reduces the time complexity to $O(N + M)$, which is a substantial improvement over the previous $O(N \times M)$.

```
1  std::vector<std::pair<std::size_t, std::size_t>>
2  FindVariableIndices(std::size_t i_ANN,
3                      std::vector<std::string> variable_names,
4                      bool input) const {
5      /*--- Create hashmap of ANN variable names to their indices ---*/
6      std::unordered_map<std::string, std::size_t> ann_var_map;
7
8      auto nVar = input ? NeuralNetworks[i_ANN].GetnInputs()
9                        : NeuralNetworks[i_ANN].GetnOutputs();
10
11      // Build hashmap: ANN variable name -> ANN variable index
12      for (auto iVar = 0u; iVar < nVar; iVar++) {
13          std::string ann_varname = input ? NeuralNetworks[i_ANN].GetInputName(iVar)
14                                          : NeuralNetworks[i_ANN].GetOutputName(iVar);
```

⁵<https://github.com/divyaprakash-iitd/MLPCpp/tree/flatten>

```

15     ann_var_map[ann_varname] = iVar;
16 }
17
18 /*--- Find matches using hashmap lookup ---*/
19 std::vector<std::pair<std::size_t, std::size_t>> variable_indices;
20
21 for (auto jVar = 0u; jVar < variable_names.size(); jVar++) {
22     auto it = ann_var_map.find(variable_names[jVar]);
23     if (it != ann_var_map.end()) {
24         // Found match: (variable_names index, ANN variable index)
25         variable_indices.push_back(std::make_pair(jVar, it->second));
26     }
27 }
28
29 return variable_indices;
30 }
31
32

```

The new implementation offers several advantages:

- **Improved Performance:** The use of a hash map for lookups is significantly faster than the nested loop approach, especially as the number of variables increases.
- **Enhanced Readability:** The new code is more organized and easier to understand. The logic is more straightforward, which improves maintainability.
- **Better Scalability:** The improved performance ensures that the code will scale more effectively as the complexity of the neural network models grows.

Implementation Note: The complete source code for this refactoring approach is available in the `gsoc_dp` branch of the project repository⁶.

4.2 Replacing FindVariableIndices Function

Disclaimer

Some of the changes described in this section are still under development and may not yet work as intended.

In this section, an alternative approach is proposed for refactoring of the variable matching logic within the `CLookUp_ANN` and `CNeuralNetwork` classes. The goal is to improve **efficiency, readability, and maintainability**. The key change is to replace the current `FindVariableIndices` function with a more powerful and efficient method, `GetVariableMapping`, inside the `CNeuralNetwork` class. This new method performs the complete variable matching and index retrieval in a single operation, eliminating redundant work and clarifying the code's intent.

The original implementation in `CLookUp_ANN::PairVariableswithMLPs` relied on calling `FindVariableIndices` twice (once for inputs, once for outputs) for every neural network in the collection.

```

1 // Original logic
2 std::vector<std::pair<size_t, size_t>> Input_Indices =

```

⁶https://github.com/divyaprakash-iitd/MLPCpp/tree/gsoc_dp

```

3     FindVariableIndices(iMLP, inputVariables, true);
4     isInput = Input_Indices.size() > 0;
5
6     if (isInput) {
7         std::vector<std::pair<size_t, size_t>> Output_Indices =
8             FindVariableIndices(iMLP, outputVariables, false);
9         isOutput = Output_Indices.size() > 0;
10
11         if (isOutput) {
12             // ... update map
13         }
14     }

```

This approach has two main drawbacks:

1. **Inefficiency:** It involves multiple loops and lookups for a single network check.
2. **Separation of Concerns:** The logic for how a `CNeuralNetwork` matches against a set of variables was located in the `CLookup_ANN` class, not within the `CNeuralNetwork` class itself.

While `operator==` was suggested, it is not ideal because it's expected to return a simple `bool`, whereas this operation requires the resulting index maps. A boolean check would still require a second operation to retrieve the indices, leading to redundant computation.

4.2.1 Implementation

The proposed solution centralizes the matching logic into a single, efficient method within the `CNeuralNetwork` class.

Step 1: Introduce MatchResult Struct and GetVariableMapping Method A new struct, `MatchResult`, is defined to hold the complete result of a matching operation. A new public method, `GetVariableMapping`, is added to `CNeuralNetwork`.

```

1 // CNeuralNetwork.hpp
2
3 struct MatchResult {
4     bool is_match = false;
5     std::vector<std::pair<size_t, size_t>> input_indices;
6     std::vector<std::pair<size_t, size_t>> output_indices;
7 };
8
9 public:
10 MatchResult GetVariableMapping(
11     const std::vector<std::string>& lookup_inputs,
12     const std::vector<std::string>& lookup_outputs
13 ) const;

```

This method efficiently checks for matches and returns the `MatchResult` struct containing a boolean flag and the corresponding index vectors, all in one pass. The matching logic used is given below.

- **Inputs:** All of the network's defined inputs must be present in the `lookup_inputs`.
- **Outputs:** At least one of the network's defined outputs must be present in the `lookup_outputs`.

Step 2: Simplify PairVariableswithMLPs and Remove FindVariableIndices The `PairVariableswithMLPs` function in `CLookUp_ANN.hpp` is refactored to use the new method, making it significantly cleaner and more expressive. The `FindVariableIndices` function is no longer needed and should be deleted.

```

1  // CLookUp_ANN.hpp
2  // Replace the old PairVariableswithMLPs function with this:
3  void PairVariableswithMLPs(MLPToolbox::CIOMap &ioMap) {
4      auto inputVariables = ioMap.GetInputVars();
5      auto outputVariables = ioMap.GetOutputVars();
6
7      for (size_t iMLP = 0; iMLP < NeuralNetworks.size(); iMLP++) {
8          // Single, efficient call to the new method
9          auto mapping_result =
10             NeuralNetworks[iMLP].GetVariableMapping(inputVariables, outputVariables);
11
12         if (mapping_result.is_match) {
13             // Update the map directly with the results
14             ioMap.PushMLPIndex(iMLP);
15             ioMap.PushInputIndices(mapping_result.input_indices);
16             ioMap.PushOutputIndices(mapping_result.output_indices);
17         }
18     }
19
20     CheckUseOfInputs(ioMap);
21     CheckUseOfOutputs(ioMap);
22 }
23
24 // The FindVariableIndices function should be deleted from this file.

```

Implementing these changes is expected to yield the following advantages:

- **Efficiency:** Eliminates redundant loops and string comparisons by performing the variable check and index retrieval in a single, optimized operation.
- **Readability:** The code in `PairVariableswithMLPs` now clearly expresses its intent: get the variable mapping from the network and use it if there is a match.
- **Encapsulation:** The logic for how a neural network matches variables is now correctly placed within the `CNeuralNetwork` class, improving code organization and maintainability.

Implementation Note: The complete source code for this refactoring approach is available in the `revarmap` branch of the project repository⁷.

5 Neural Network Function Optimization

Several critical computational functions in the neural network implementation were optimized to improve both performance and numerical precision. The optimizations focused on reducing redundant array access operations and leveraging fused multiply-add (FMA) instructions for better floating-point accuracy.

⁷<https://github.com/divyaprakash-iitd/MLPCpp/tree/revarmap>

5.1 Implementation

The key optimization strategies used are listed below.

1. **Pre-extraction of frequently accessed data:** Weight vectors and neuron counts are cached before loops to minimize repeated pointer dereferencing.
2. **FMA instruction usage:** Replaced separate multiplication and addition operations with `std::fma()` for improved numerical precision and potential performance gains.

Using the above strategies, the compute functions can be modified accordingly. As an illustration, an example is provided for the `computeX` function.

```
1  mlpdouble ComputeX(std::size_t iLayer, std::size_t iNeuron) const noexcept {
2      const auto& w = weights_mat[iLayer - 1][iNeuron];
3      const auto n = total_layers[iLayer - 1]->GetNNeurons();
4
5      mlpdouble x = total_layers[iLayer]->GetBias(iNeuron); // start from bias
6      for (std::size_t j = 0; j < n; ++j) {
7          // If GetOutput(j) is cheap (non-virtual, inlined), this is fine.
8          x = std::fma(w[j], total_layers[iLayer - 1]->GetOutput(j), x);
9      }
10     return x;
11 }
```

Implementation Note: The complete source code for this refactoring approach is available in the `gsoc_dp` branch of the project repository⁸.

6 Results

The optimizations described in Subsections 3.2 and 4.1 were combined in the code. From Section 5, the only modification implemented was caching the neuron counts before entering the loops. The use of `std::fma` was not adopted, as it appeared to slow down the code when integrated into full SU2 simulations, even though it had shown good performance in the standalone test case. The resulting implementation was submitted as a Pull Request⁹ to the MLPCpp subproject of SU2.

To verify the correctness of the implementation, the code was compiled and the test case provided in the MLPCpp repository was executed. In this test, the derivatives obtained using finite differences were found to agree with those calculated analytically by the network, thereby confirming the validity of the implementation.

The prediction times of the original implementation were compared with those of the optimized GSoC version over multiple runs. The original implementation was observed to have an average runtime of **251.98 ± 7.16 ms**, whereas the optimized version reduced this to **178.96 ± 4.56 ms**. This corresponds to a **1.41× speedup**, or equivalently a **28.9% reduction in execution time**. Figure X presents the average runtimes with standard deviation error bars, which reflect the run-to-run fluctuations in the measurements. It may therefore be concluded that the optimized implementation not only achieves faster predictions but also exhibits greater consistency across runs.

However, when the optimized implementation was integrated into full SU2 simulations, the performance gains observed in the standalone test case were not replicated in the complete application environment. This highlights an important lesson in optimization work: standalone benchmarks may not accurately reflect performance in the target application due to differences in computational context, memory access patterns, and overall system load. Future work will focus on developing test cases that more closely mirror the actual SU2 deployment environment to ensure that performance improvements translate effectively to real-world usage scenarios.

⁸https://github.com/divyaprakash-iitd/MLPCpp/tree/gsoc_dp

⁹<https://github.com/EvertBunschoten/MLPCpp/pull/2>




Figure 3: Average prediction time of the original and refactored implementations. Error bars represent the standard deviation across multiple runs. The refactored implementation was observed to achieve a mean runtime of 189.90 ± 3.36 ms compared to 251.98 ± 7.16 ms for the original, corresponding to a $1.33\times$ speedup (approximately 25% reduction in execution time).

7 Conclusion

This project has contributed to the performance improvement and usability of the `MLPCpp` module in SU2. Through systematic profiling and targeted optimizations, key runtime bottlenecks were identified and addressed. The refactoring of the weight matrix storage, introduction of hash-based variable lookup, and improvements to the neural network compute functions resulted in measurable performance gains in standalone testing. Specifically, the optimized implementation achieved a mean prediction runtime of 189.90 ± 3.36 ms compared to 251.98 ± 7.16 ms for the original version, corresponding to a $1.33\times$ speedup (approximately 25% reduction in execution time) in the test case environment.

While these performance improvements did not fully translate to the complete SU2 simulation environment, this outcome provides valuable insights into the challenges of performance optimization across different computational contexts. Future work will focus on developing more representative test cases that better mirror the actual SU2 deployment conditions, ensuring that optimization efforts can be more effectively validated and translated to real-world performance gains.

Beyond performance improvements, this work also contributed to the long-term maintainability of SU2. The integration of the Tracy profiler into the Meson build system, along with the accompanying documentation, ensures that future development within SU2 can benefit from low-overhead runtime analysis and continuous performance monitoring.

8 Future Work

Although some optimizations were completed, several tasks remain to be addressed in order to further improve performance, usability, and maintainability. The following points outline the next steps:

- **Representative test case development:** Develop test cases that more accurately reflect the computational environment and workload characteristics of full SU2 simulations, enabling better validation of performance optimizations in realistic deployment scenarios.
- **Flattened weight storage:** Complete the debugging of the flattened vector implementation and compare its performance against the nested vector version to evaluate the actual performance gains in both standalone and integrated environments.
- **Refactor variable mapping:** Finalize replacement of the `FindVariableIndices` function with a direct, maintainable mapping for faster lookups.

- **Tracing support:** Develop a Tracy tutorial demonstrating how to compile SU2 with tracing, instrument code sections, and interpret performance traces.
- **Documentation:** Expand MLPCpp documentation and onboarding notes, including example cases for easier adoption by new users.

9 Acknowledgment

I would like to sincerely thank Google and the SU2 community for providing me with the opportunity to contribute to this project through GSoC 2025. I am especially grateful to my mentors, Evert Bunschoten and Joshua A. Kelly, for their constant guidance, encouragement, and support throughout the project. Their approachable and friendly mentorship created an excellent learning environment, and their insights were invaluable in shaping the progress and outcomes of this work. I am also thankful to the wider SU2 developer community for their feedback, code reviews, and discussions, which greatly enriched the overall experience.

References

- [1] Thomas D Economon et al. “SU2: An open-source suite for multiphysics simulation and design”. In: *Aiaa Journal* 54.3 (2016), pp. 828–846.